

# Naïve Bayes II

Semester 1, 2023


Kris Ehinger

# Outline





- Naïve Bayes in practice
- Continuous data
- Converting data types
- Naïve Bayes with continuous data

# Naïve Bayes example

- Simple shapes dataset:
  - Tall or Wide
  - Oval or Rectangle
  - Red, Yellow, Green, or Blue




Tall, Rectangle, Red



Wide, Rectangle, Green

# Naïve Bayes example



# Naïve Bayes in practice

- Practical considerations
  - How to handle zero probabilities? ← Previous lecture
  - How to handle missing values?
  - How to handle small probabilities (underflow errors)
  - What if attributes aren't conditionally independent?

# Missing values

- What if an instance is missing some attribute?
- Missing values at test can simply be ignored – compute the likelihood of each class from the non-missing values
- Missing values in training can also be ignored – don't include them in the attribute-class counts,<sup>N</sup> and the probabilities will be based on the non-missing values

# Avoiding underflow

- Multiplying a lot of numbers (0,1] can lead to underflow – numbers smaller than the computer can represent
- Common workaround: **log transformation**
  - take the log() of each probability and sum instead of multiplying

$$\begin{aligned}\hat{c} &= \arg \max_{c_j \in C} P(c_j) \prod_i P(x_i | c_j) \\ &= \arg \max_{c_j \in C} \left[ \log(P(c_j)) + \sum_i \log(P(x_i | c_j)) \right]\end{aligned}$$

*Handwritten annotations:* "multiplying" with an arrow pointing to the product symbol  $\prod_i$ ; "take log" with an arrow pointing to the  $\log$  function in the transformed equation.

# Independence assumption

- Independence assumption is usually wrong
- But Naïve Bayes usually works anyway. Why?
  - We don't need a perfect estimate of  $P(c|T)$  for every class – we just need to know which class is most likely
  - Ignoring the fact that some attributes are correlated tends to make all the class probabilities higher, but doesn't typically change their rank
  - Naïve Bayes is also robust to small errors in estimating  $P(x_i|c_j)$  – these can change class probabilities but typically don't change class rank



# Naïve Bayes

- Strengths:
  - Simple to build, fast
  - Computations scale well to high-dimensional datasets (1000s of attributes)
  - Explainable – generally easy to understand why the model makes the decision it does
- Weaknesses:
  - Inaccurate when there are <sup>likelihood</sup> many missing  $P(x_i | c_j)$  values
  - Conditional independence assumption becomes problematic for complex systems <sub>not all dataset</sub>

# Continuous data

# Continuous attributes

- Naïve Bayes (as discussed last lecture) assumes nominal data
  - What happens if we have continuous data?

# Continuous attributes

- How to compute probabilities  $P(x_i | c_j)$ ?

Wind	Temp	Rain?
north	32.1	no
east	18.4	yes
east	19.5	no
north	23.5	no
north	20.3	yes
east	19.7	yes

$$P(\text{temp} = 32.1 | \text{Rain} = \text{no}) = 1/3$$

$$P(\text{temp} = 19.5 | \text{Rain} = \text{no}) = 1/3$$

$$P(\text{temp} = 23.5 | \text{Rain} = \text{no}) = 1/3$$

$$P(\text{temp} = 18.4 | \text{Rain} = \text{yes}) = 1/3$$

$$P(\text{temp} = 20.3 | \text{Rain} = \text{yes}) = 1/3$$

$$P(\text{temp} = 19.7 | \text{Rain} = \text{yes}) = 1/3$$

This doesn't look right...

# Converting data types

# Data types

- The input to a machine learning system consists of instances, which have:
  - Attributes
  - Class labels (if supervised)
- Attributes and class labels can be:
  - Nominal / categorical
  - Ordinal
  - Continuous / numeric

# Attribute types

- Machine learning algorithms typically assume attributes have a particular data type
- Algorithms that assume nominal attributes:
  - Naïve Bayes (as described in last lecture)
  - Decision trees
- Algorithms that assume numeric attributes:
  - Support vector machines (SVM)
  - Perceptron, neural network


# Attribute types

- What if we have attributes of the wrong type for a given model?
- Options:
  - Discard those attributes
  - Convert the attributes to match the model
  - Change the model assumptions to match the data



# Nominal -> numeric


- Option 1: Convert category names to numbers
  - Attribute: mode of transport
  - Nominal: “bus,” “tram,” “car,” “bike,” “walk” *not ordinal*
  - Numeric: 0, 1, 2, 3, 4
- Problem: creates an artificial ordering when no order exists, makes some categories seem more/less similar to each other



# Nominal -> numeric

- Option 2: One-hot encoding
  - Attribute with  $m$  possible values ->  $m$  boolean attributes
  - “bus” = [1, 0, 0, 0, 0]
  - “tram” = [0, 1, 0, 0, 0]
  - “car” = [0, 0, 1, 0, 0]
  - “bike” = [0, 0, 0, 1, 0]
  - “walk” = [0, 0, 0, 0, 1]
- Best way to represent nominal values in a continuous space, but increases dimensionality of the space  
*1 attribute with  $m$  values  $\Rightarrow m$  attributes*

# Visualisation of option 2



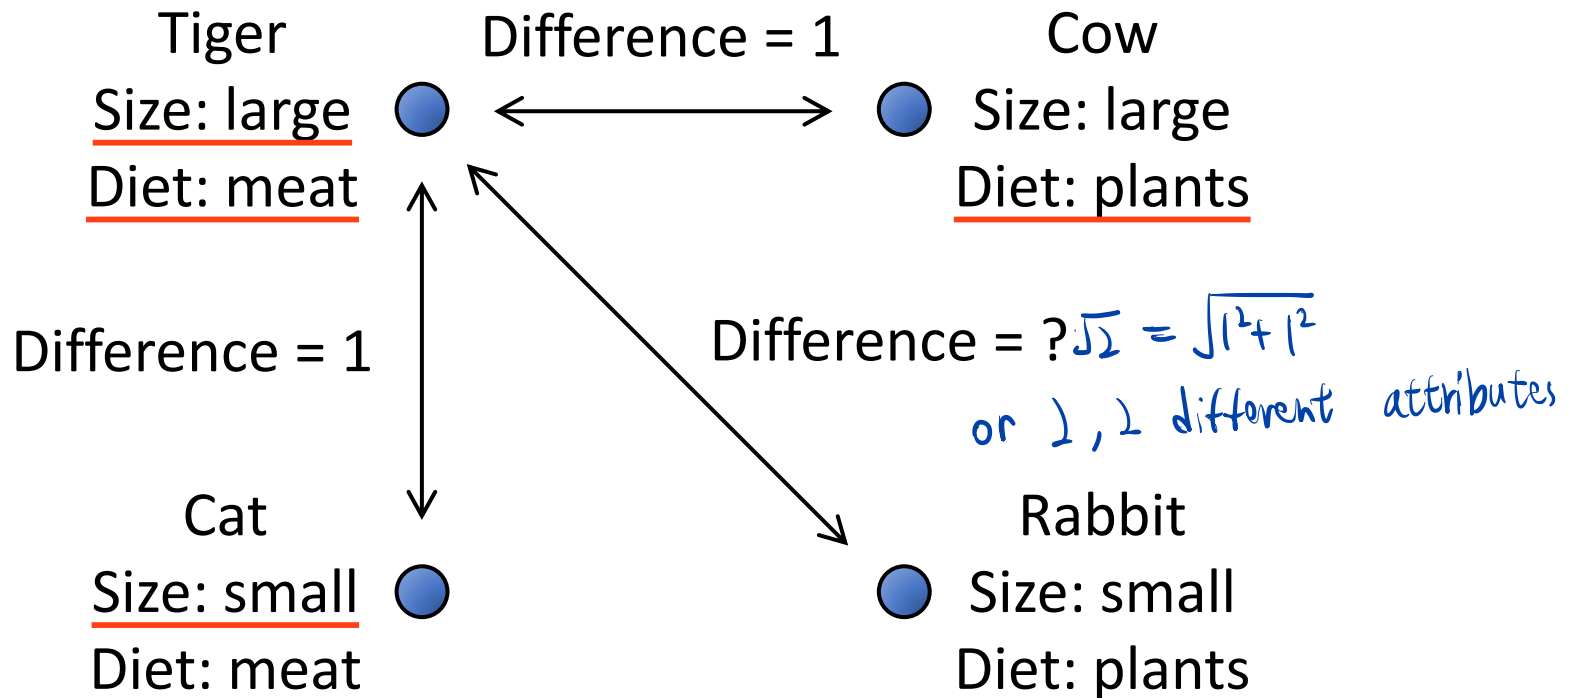
## Better encoding:

All categories equal distance apart; any subset is equally easy to group

# Computing distances

- How to compute distances between nominal attributes?
- Consider these boolean attributes for an animal classifier:
  - Size: large/small
  - Diet: meat/plants

# Computing distances (difference)



# Computing distances

- How to compute distances between nominal attributes?

- Euclidean distance

- If A and B differ on N attributes,  $dist_E(A, B) = \sqrt{N}$

$$dist_E(A, B) = \sqrt{\sum_i (a_i - b_i)^2}$$

- Hamming distance

- If A and B differ on N attributes,  $dist_H(A, B) = N$

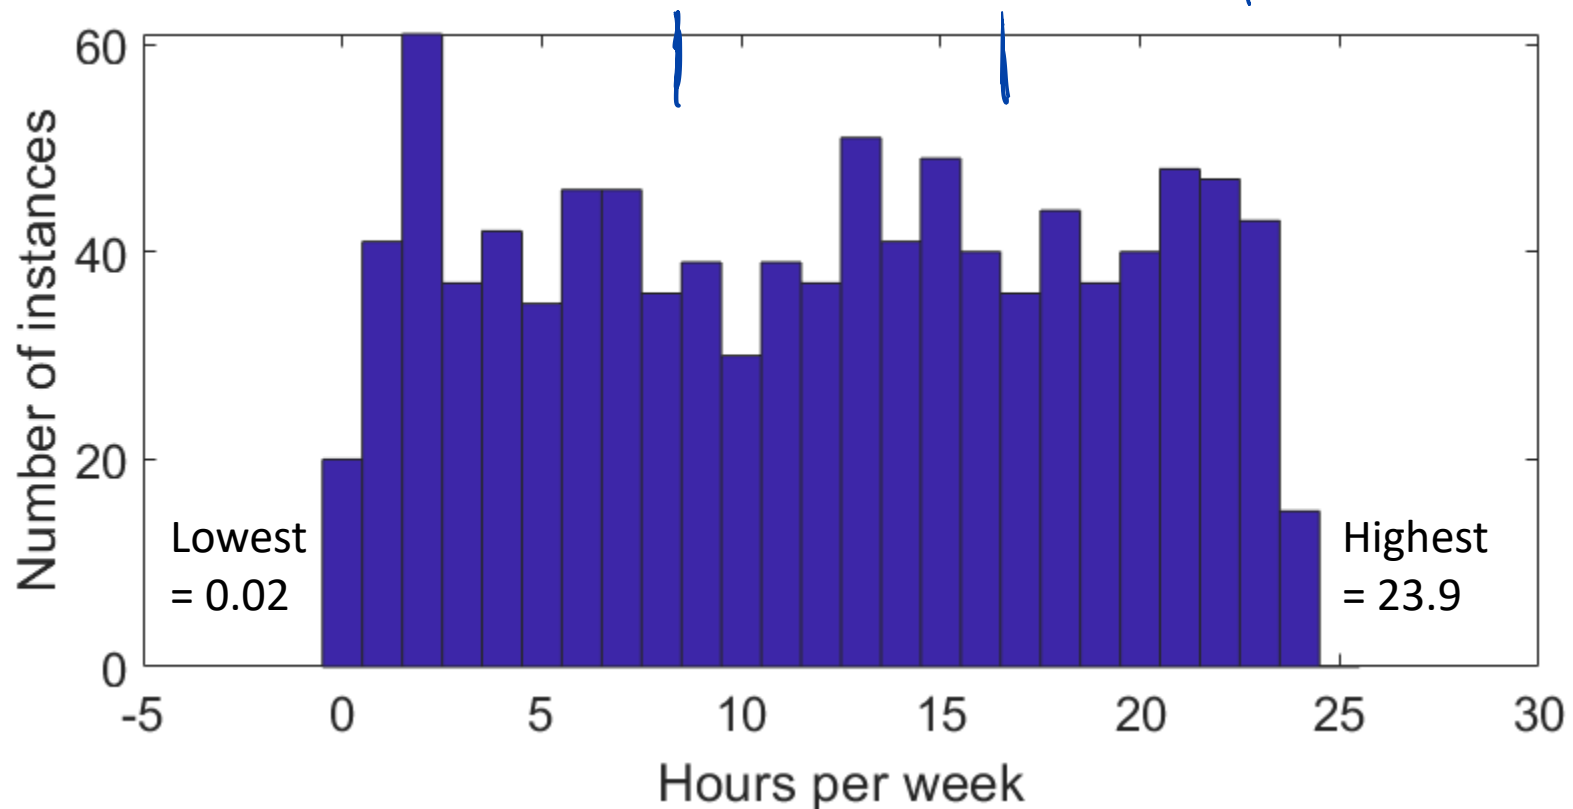
$$dist_H(A, B) = \sum_i \begin{cases} 0, & a_i == b_i \\ 1, & otherwise \end{cases}$$

# Numeric -> nominal

- **Discretisation** is the translation of continuous numeric attributes to discrete nominal attributes
  - Example: map temperatures to “hot,” “mild,” “cool”
- Discretisation is generally a two-step process:
  - Decide how many nominal values (= intervals) onto which you will map the numeric values *# bins*
  - Decide where to place the boundaries for these intervals

# Discretisation example 1

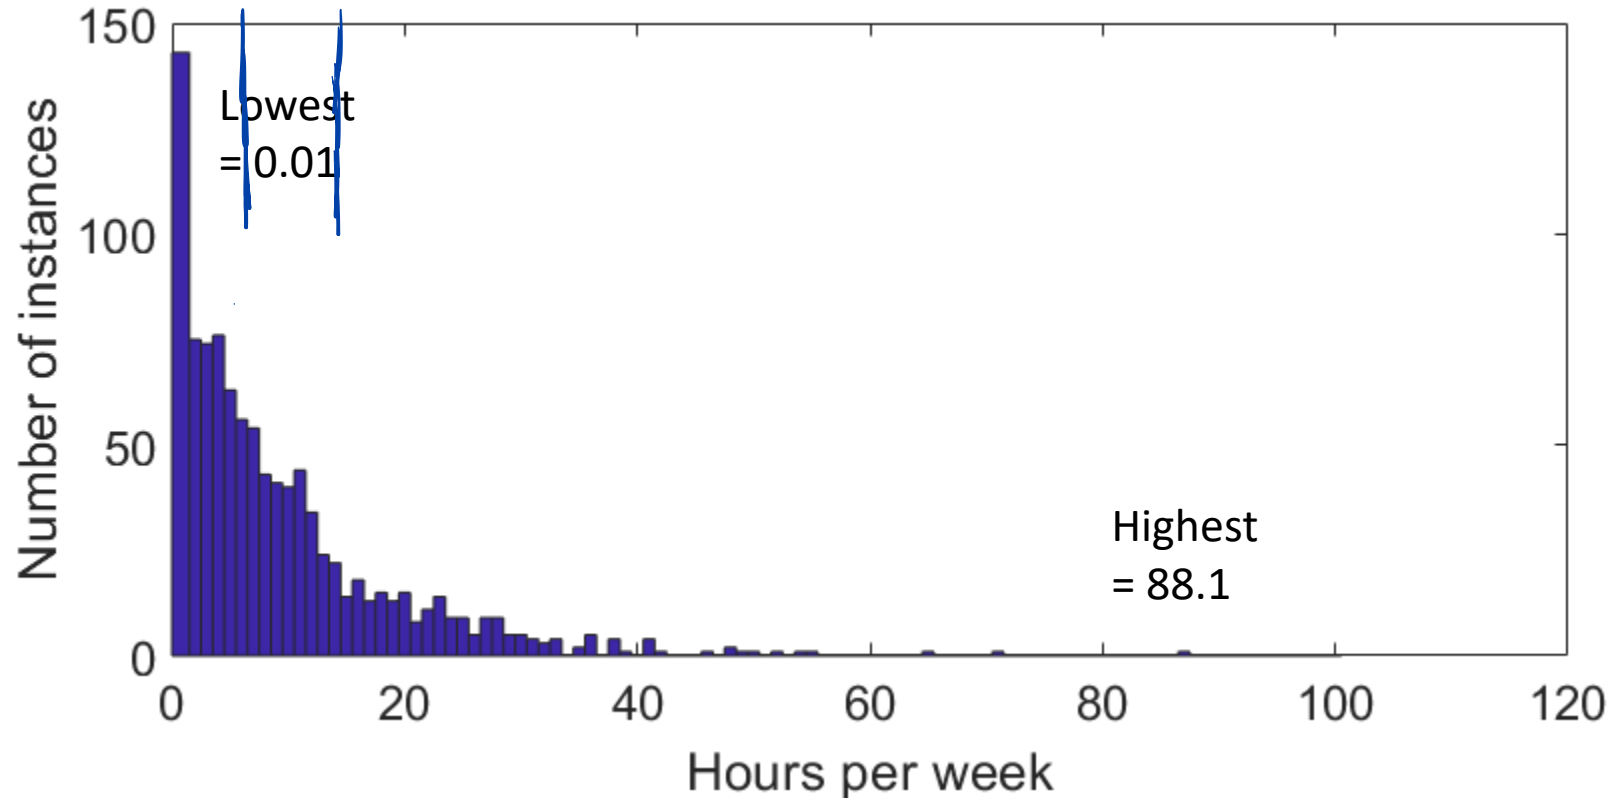
You measured app usage (hours/week) from 1000 users. How would you recode this attribute into 3 levels: low, medium, high? *equal-width*






# Discretisation example 2

You measured app usage (hours/week) from 1000 users. How would you recode this attribute into 3 levels: low, medium, high? *equal frequency*



# Discretisation example 3

You measured app usage (hours/week) from 1000 users. How would you recode this attribute into 3 levels: low, medium, high?




# Discretisation options

- <sup>e.g. 1</sup> **Equal-width** discretisation – find min/max of range,  
partition into n bins of width (max-min)/n
- <sup>e.g. 2</sup> **Equal-frequency** discretisation – sort values, find  
breakpoints that produce n bins with  
(approximately) equal numbers of items
- Disadvantages:
  - Arbitrary group boundaries
  - Equal-width is sensitive to outliers, equal-frequency is  
sensitive to sample bias
  - User must choose n


# Discretisation options

- Use some clustering method to discover natural breakpoints within your data (e.g., k-means clustering)
- Disadvantages
  - More complicated than equal-width or -frequency
  - If data doesn't have natural "groups," k-means result is the same as equal-width discretisation
  - Sensitive to outliers
  - User must choose n


# Discretisation: K-means clustering



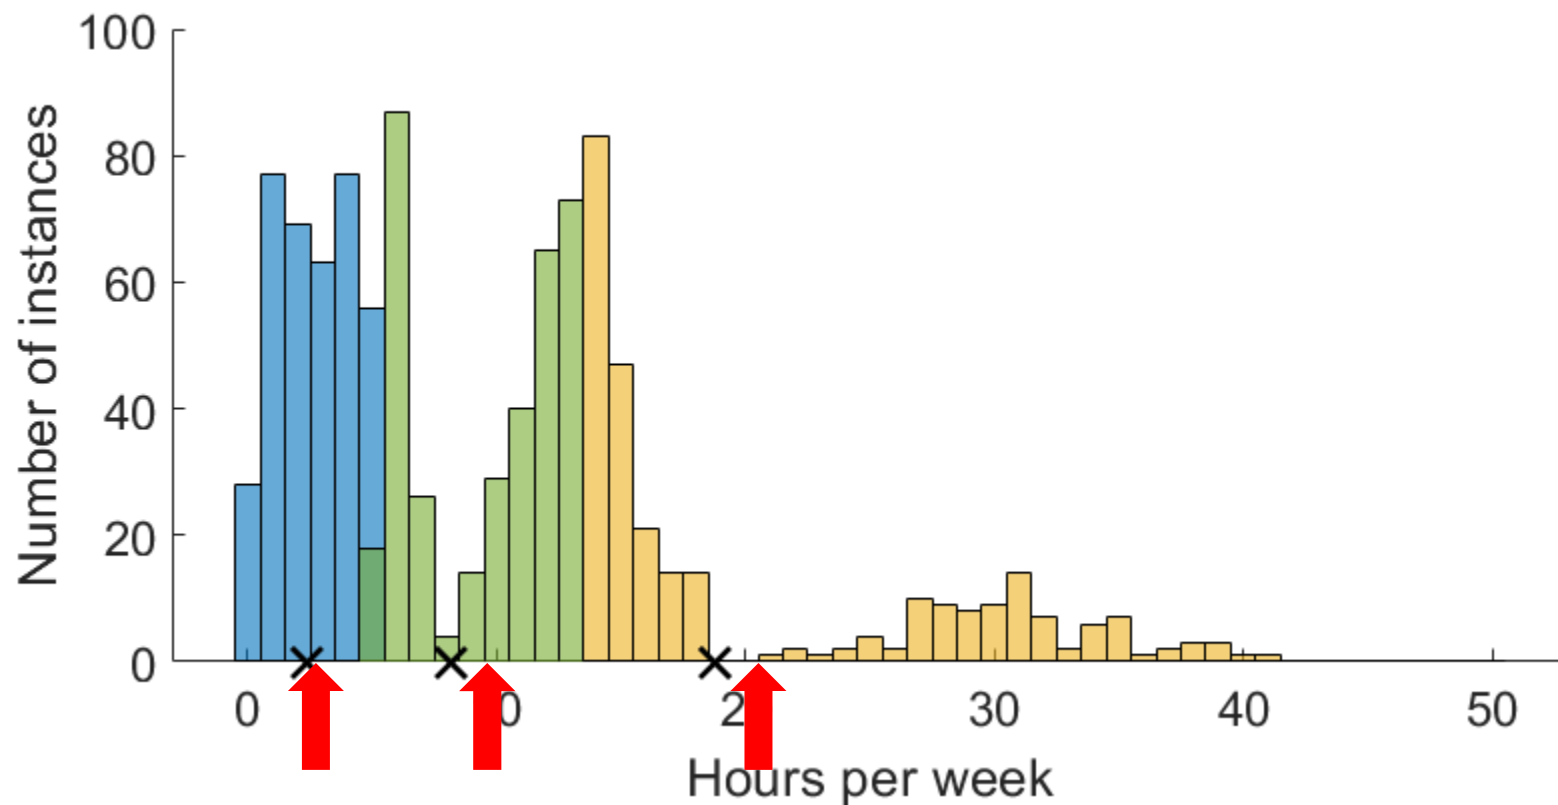
# Discretisation: K-means clustering



# Discretisation: K-means clustering




# Discretisation: K-means clustering







# Discretisation: K-means clustering




# Discretisation: K-means clustering



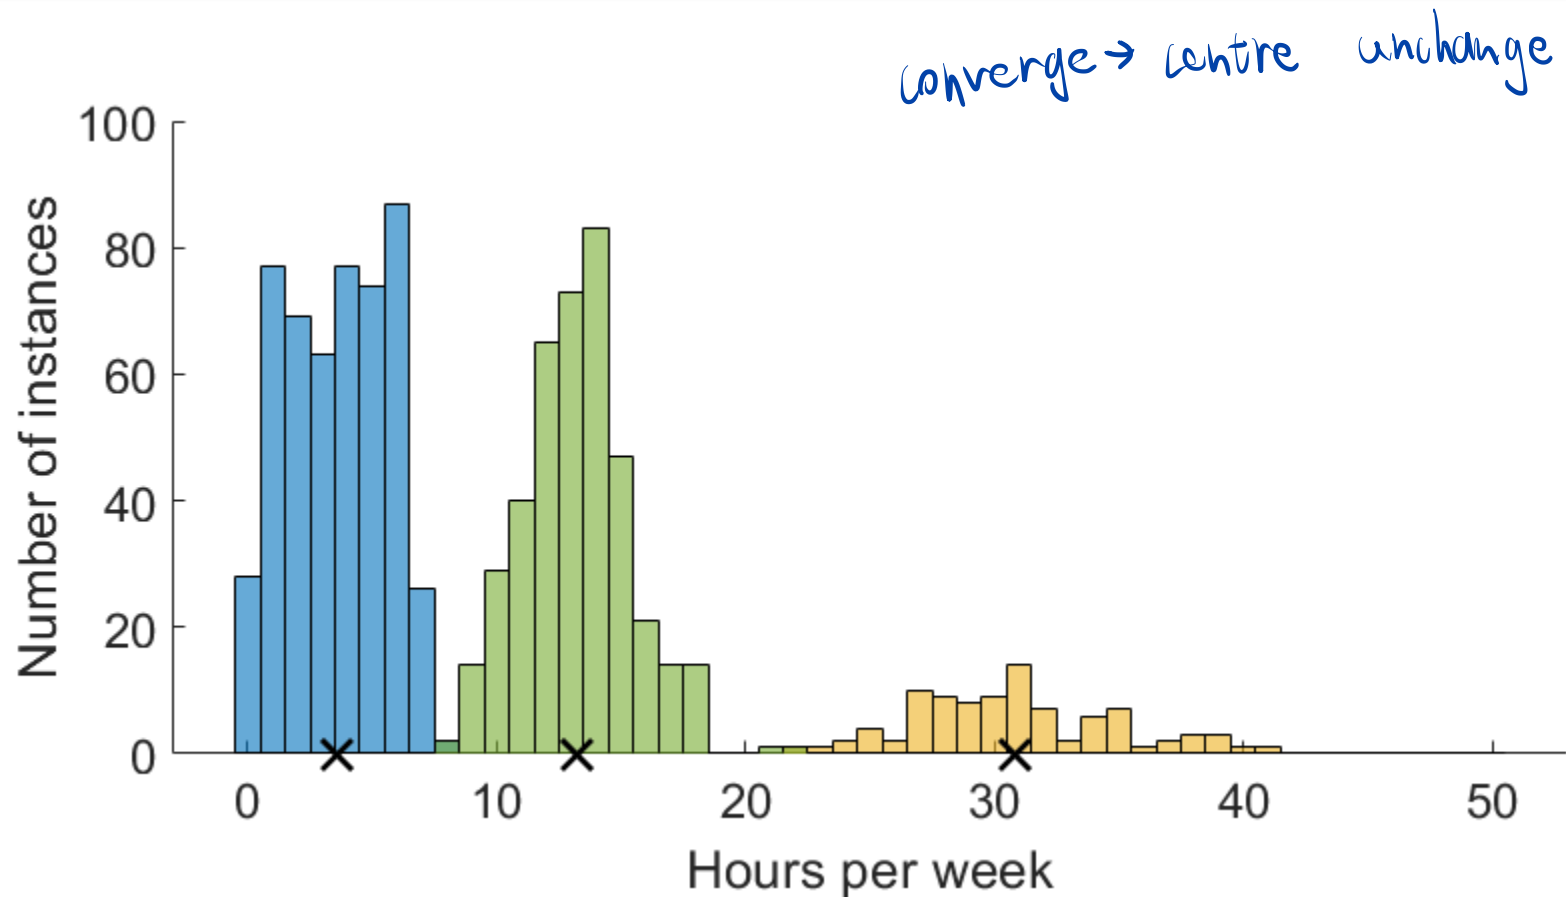
# Discretisation: K-means clustering



# Discretisation: K-means clustering



# Discretisation: K-means clustering



# Supervised discretisation

- Group values into class-contiguous intervals

- Sort values, and identify breakpoints in class membership

temp → 1 | 2 | 4 5 7 | 8 9 | 9 11  
 rainy not rainy rainy → label

- Reposition breakpoints if the numeric value is the same

1 | 2 | 4 5 7 | 8 9 9 | 11

- Set the breakpoints midway between the neighbouring values

1 | 2 | 4 5 7 | 8 9 9 | 11  
 1.5 3 7.5 10

5 bins

# Supervised discretisation

- Supervised discretisation may help you find “groups” that are most relevant for your classification task *↳ to the label*
- Disadvantages
  - Arbitrary group boundaries
  - Arbitrary number of groups (tends to produce too many groups)
  - Tends to **overfit** training set

# Discretisation summary

- Various options; each has advantages and disadvantages
- Discretisation means throwing out some information
  - This can help simplify data to make it easier for the model to learn
  - But you can potentially lose details that would have been useful



# Naïve Bayes with continuous data

# Naïve Bayes

- In naïve Bayes, we choose the class  $c_j$  that maximizes:

$$\text{posterior} \quad P(c_j) \prod_i P(x_i | c_j)$$


Prior probability of class  $c_j$       Likelihood of feature  $x_i$  in class  $c_j$

- $P(x_i | c_j)$  can be computed differently for different data types:
  - If  $x_i$  is a nominal attribute with multiple levels, count number of times each level occurs in class  $c_j$
  - If  $x_i$  is a numeric attribute, compute its probability density function

pdf

# Probability density function

- A popular pdf for continuous data is the **Gaussian distribution** (or normal distribution)
- The probability of observing value x from a variable with mean (expected value)  $\mu$  and standard deviation  $\sigma$ :



$$f(x) = \phi_{\mu,\sigma}(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$

$\phi_\sigma$  is a Gaussian distribution with mean = 0 and standard deviation =  $\sigma$

- Important properties of this distribution:
  - Symmetric about the mean  $\mu$
  - Area under the curve sums to 1

# Probability density function

- Where do the mean  $\mu$  and standard deviation  $\sigma$  come from?
- These are descriptive statistics that can be estimated from our data

Mean of N samples  
of an attribute X:

$$\mu_x = \sum_{i=1}^N \frac{x_i}{N}$$

Standard deviation of N  
samples of an attribute X:

$$\sigma_x = \sqrt{\frac{\sum_{i=1}^N (x_i - \mu)^2}{N - 1}}$$

# Why Gaussian?

- In practice, the normal distribution is a reasonable approximation for many events
  - Including binomial, Poisson, chi-square, and Student's t-distributions
  - E.g., height, weight, shoe size, reaction time ...
- Easy to compute mean and standard deviation from data
- Even if the data isn't quite normally distributed, assuming a normal distribution often works well enough

# Example: Gaussian naïve Bayes

- Compute probability distribution for each class:

Wind	Temp	Rain?
north	32.1	no
east	18.4	yes
east	19.5	no
north	23.5	no
north	20.3	yes
east	19.7	yes

When rain=no, temp is  
[32.1, 19.5, 23.5]  $N=\}$

Mean:

$$\mu_{no} = (32.1 + 19.5 + 23.5)/3 = 25.0$$

Standard deviation:

$$\sigma_{no} = \sqrt{\frac{(32.1 - 25.0)^2 + (19.5 - 25.0)^2 + (23.5 - 25.0)^2}{3 - 1}} = 6.4$$

# Example: Gaussian naïve Bayes

- Compute probability distribution for each class:

Wind	Temp	Rain?
north	32.1	no
east	18.4	<b>yes</b>
east	19.5	no
north	23.5	no
north	20.3	<b>yes</b>
east	19.7	<b>yes</b>

When rain=yes, temp is  
[18.4, 20.3, 19.7]

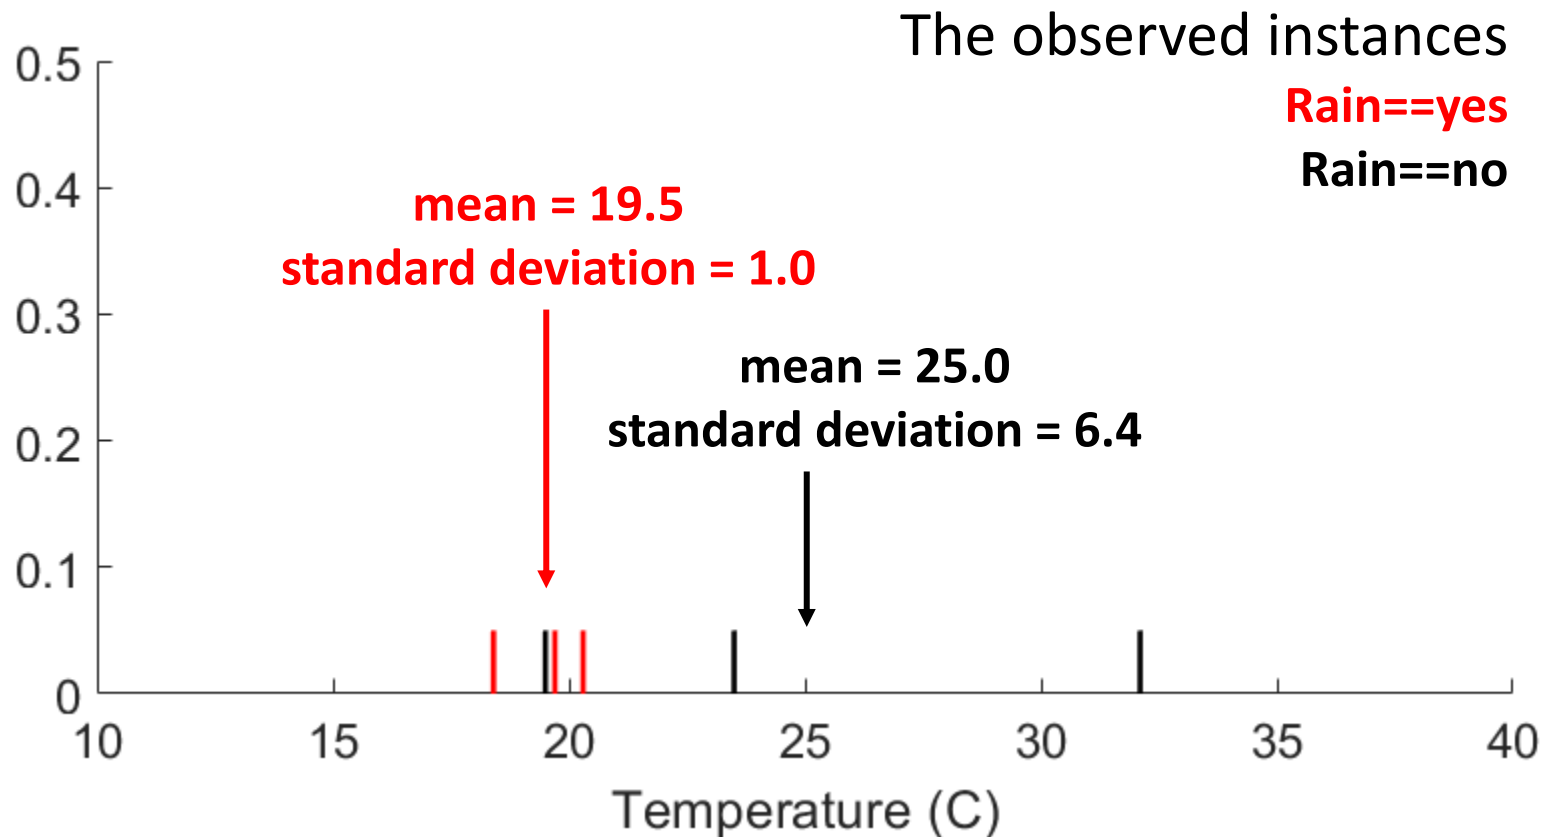
Mean:

$$\mu_{yes} = (18.4 + 20.3 + 19.7)/3 = 19.5$$

Standard deviation:

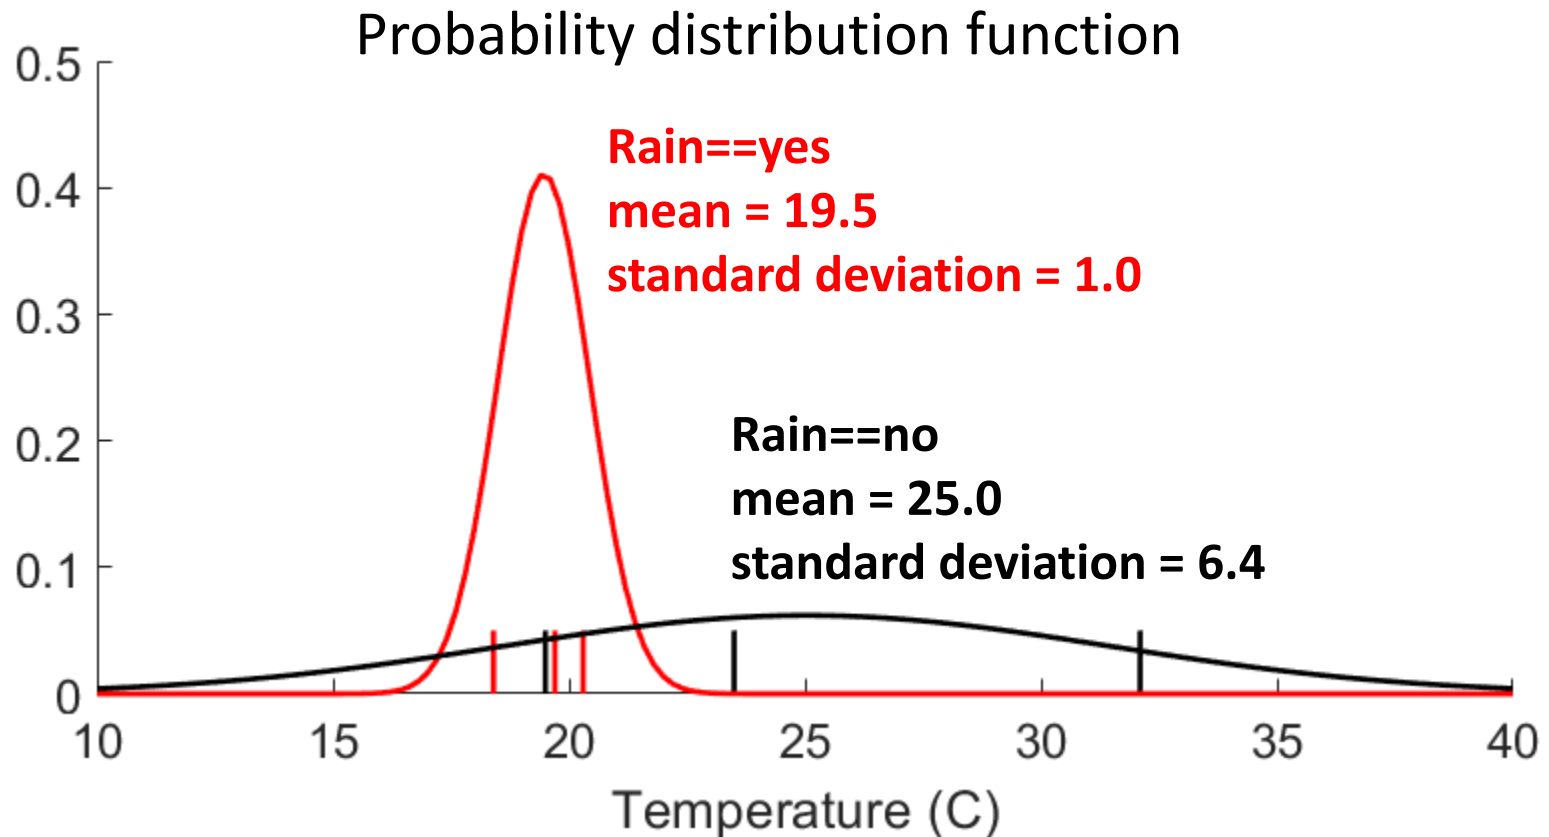
$$\sigma_{yes} = \sqrt{\frac{(18.4 - 19.5)^2 + (20.3 - 19.5)^2 + (19.7 - 19.5)^2}{3 - 1}} = 1.0$$

# Example: Gaussian naïve Bayes

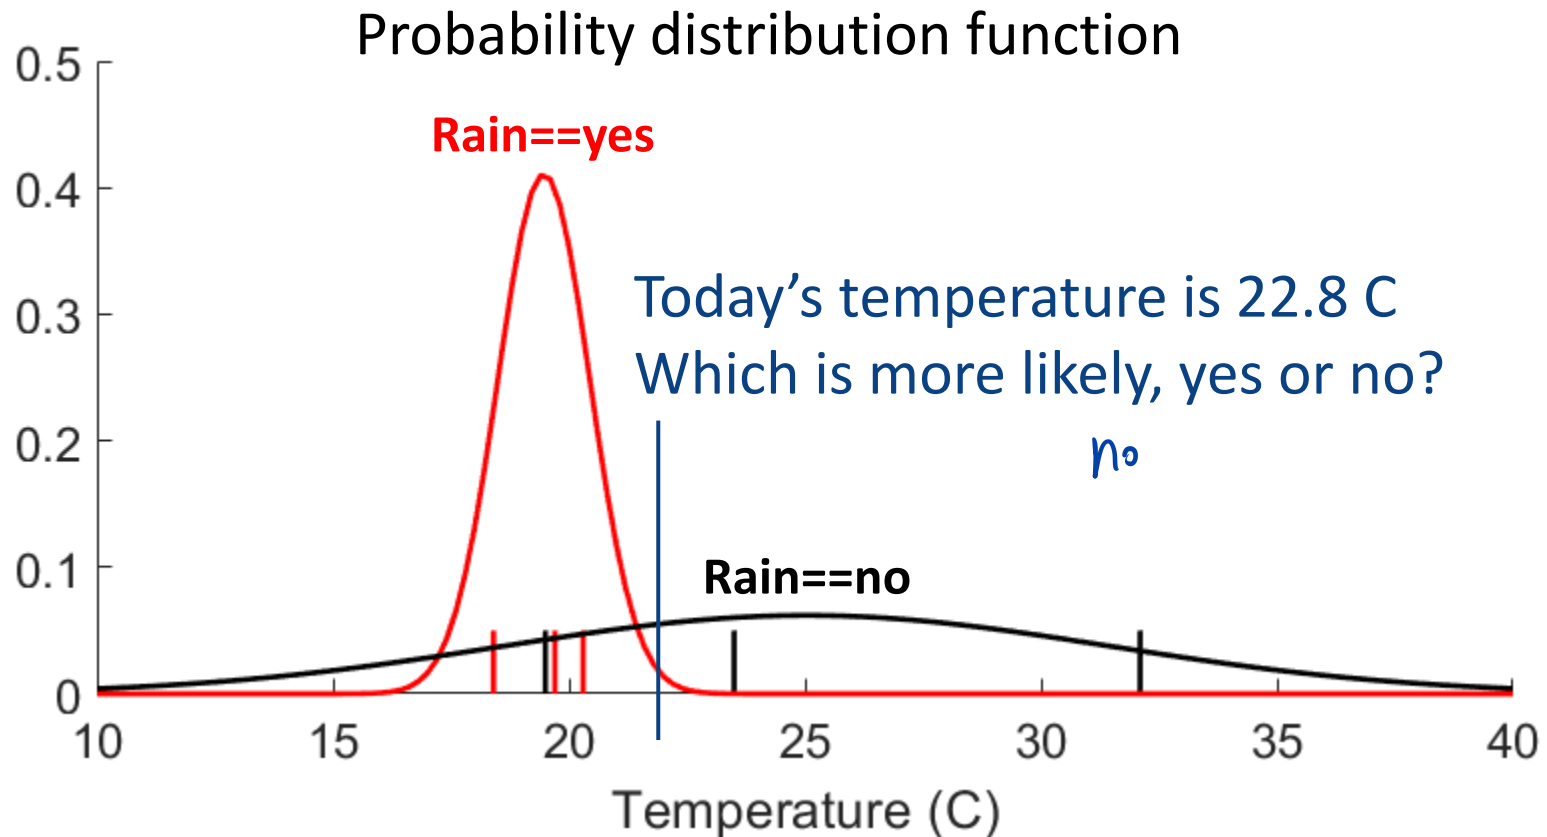




# Example: Gaussian naïve Bayes



# Example: Gaussian naïve Bayes



# Example: Gaussian naïve Bayes

- Compute the probability of the test instance by substituting the  $\mu$  and  $\sigma$  for each class:

$$P(\text{rain} = \text{no} | \text{temp} = 22.8) \stackrel{(\sim)}{=} P(\text{rain} = \text{no})P(\text{temp} = 22.8 | \text{rain} = \text{no})$$

$$P(\text{rain} = \text{yes} | \text{temp} = 22.8) \stackrel{(\sim)}{=} P(\text{rain} = \text{yes})P(\text{temp} = 22.8 | \text{rain} = \text{yes})$$

# Example: Gaussian naïve Bayes

- Compute the probability of the test instance by substituting the  $\mu$  and  $\sigma$  for each class:

$$P(\text{rain} = \text{no} | \text{temp} = 22.8) \approx P(\text{rain} = \text{no}) \phi_{\mu_{\text{no}}, \sigma_{\text{no}}}(22.8)$$


$$= (0.5) \frac{1}{6.4\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{22.8-25.0}{6.4}\right)^2} = \mathbf{0.0292}$$

$$P(\text{rain} = \text{yes} | \text{temp} = 22.8) \approx P(\text{rain} = \text{yes}) \phi_{\mu_{\text{yes}}, \sigma_{\text{yes}}}(22.8)$$

$$= (0.5) \frac{1}{1.0\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{22.8-19.5}{1.0}\right)^2} = \mathbf{0.0006}$$

take by better

# Explaining naïve Bayes



# Naïve Bayes for mixed data types

- In naïve Bayes, we choose the class  $c_j$  that maximizes:

$$P(c_j) \prod_i P(x_i | c_j)$$

Prior probability of class  $c_j$       Likelihood of feature  $x_i$  in class  $c_j$

- How you compute  $P(x_i | c_j)$  depends on the data type of  $x_i$  and its probability distribution or density function

# Types of naïve Bayes

- **Multivariate**: attributes are nominal and can take any of a fixed number of values
- Binomial (or Bernoulli): attributes are binary *model assumptions:*
  - Special case of multivariate
- Multinomial: attributes are natural numbers corresponding to frequency *model assumptions:*
  - Probability  $P(a_k = m | c_j) \approx P(a_k = 1 | c_j)^m / (m!)$
- **Gaussian**: attributes are numeric, and we can assume they come from a Gaussian distribution *model assumptions:*
- Kernel density estimation: attributes are numeric, and come from an arbitrary distribution *learn the distribution from data*

# Naïve Bayes summary

- Naïve Bayes works with various data types
  - Just need to change how you compute  $P(x_i | c_j)$  for each attribute
- Mix numeric and nominal attributes in one model – no need to convert data types
- Optimal combination of attributes
  - Assuming conditional independence, and correct estimates of the probability distributions